



Section 4: Retrieval-based LMs: Training

Retrieval-based LMs




Training retrieval-based LMs



Why is training challenging?




Datastore



Why is training challenging?




Datastore




Too large! Expensive to update index during training!

Challenges of updating retrieval models




Challenges of updating retrieval models



7

Challenges of updating retrieval models




Re-indexing will be very expensive!

Why is training challenging?



Datastore



Training methods for retrieval-based LMs

- Independent training
- Sequential training
- Joint training w/ asynchronous index update
- Joint training w/ in-batch approximation

Training methods for retrieval-based LMs

- **Independent training**
 - Sequential training
 - Joint training w/ asynchronous index update
 - Joint training w/ in-batch approximation


Independent training

Retrieval models and language models are trained **independently**

- Training language models



- Training retrieval models




Independent training

Retrieval models and language models are trained **independently**

- Training language models



- Training retrieval models



Training language models



Minimize $-\log P_{\text{LM}}(y | x)$


Training language models




Minimize $-\log P_{\text{LM}}(y | x)$



GPT



PaLM



LLaMA



GPT-J

.....


Independent training

Retrieval models and language models are trained **independently**

- Training language models



- Training retrieval models



Sparse retrieval models: TF-IDF / BM25

In 1997, Apple merged with NeXT,
and Steve Jobs became CEO of ...

Jobs returned to Apple as CEO
after the company's acquisition ...

Text chunks

[0, 0, 0.4, 0, 0.8, 0.7, ...]

[0, 1.2, 0.4, 0, 0.8, 0, ...]


Sparse vectors




Lexical overlap

No training needed!


Dense retrieval models: DPR (Karpukhin et al. 2020)



Dense retriever: Inference




Dense retriever: Inference




How to train dense retrieval models?

Training dense retrieval models: DPR




Training dense retrieval models: DPR



$$L(q, p^+, p_1^-, p_2^-, \dots, p_n^-) = -\log \frac{\exp(\text{sim}(q, p^+))}{\exp(\text{sim}(q, p^+)) + \sum_{j=1}^n \exp(\text{sim}(q, p_j^-))}$$


Training dense retrieval models: DPR



$$L(q, p^+, p_1^-, p_2^-, \dots, p_n^-) = -\log \frac{\exp(\text{sim}(q, p^+))}{\exp(\text{sim}(q, p^+)) + \sum_{j=1}^n \exp(\text{sim}(q, p_j^-))}$$


Contrastive loss

Training dense retrieval models: DPR




$$L(q, p^+, p_1^-, p_2^-, \dots, p_n^-) = -\log \frac{\exp(\text{sim}(q, p^+))}{\exp(\text{sim}(q, p^+)) + \sum_{j=1}^n \exp(\text{sim}(q, p_j^-))}$$

Contrastive loss




Training dense retrieval models: DPR



$$L(q, p^+, p_1^-, p_2^-, \dots, p_n^-) = -\log \frac{\exp(\text{sim}(q, p^+))}{\exp(\text{sim}(q, p^+)) + \sum_{j=1}^n \exp(\text{sim}(q, p_j^-))}$$

Positive passage

Training dense retrieval models: DPR



Negative passages
Too expensive to consider all negatives!

$$L(q, p^+, p_1^-, p_2^-, \dots, p_n^-) = -\log \frac{\exp(\text{sim}(q, p^+))}{\exp(\text{sim}(q, p^+)) + \sum_{j=1}^n \exp(\text{sim}(q, p_j^-))}$$


Training with “in-batch” negatives

$$L(q, p^+, p_1^-, p_2^-, \dots, p_n^-) = -\log \frac{\exp(\text{sim}(q, p^+))}{\exp(\text{sim}(q, p^+)) + \sum_{j=1}^n \exp(\text{sim}(q, p_j^-))}$$

Training batch	
q_1	p_1^+ ...It was incorporated by Jobs and Wozniak as Apple Computer, Inc. in 1977. ...
q_2	p_2^+ 12-year-old Spanish football club Real Madrid is undoubtedly the best football club Spain has ever...
...	...
q_n	p_n^+ Thomas Umunnakwe Aguiyi-Ironsi seized power during the ensuing chaos after the 15 January ...

Training with “in-batch” negatives


$$L(q, p^+, p_1^-, p_2^-, \dots, p_n^-) = -\log \frac{\exp(\text{sim}(q, p^+))}{\exp(\text{sim}(q, p^+)) + \sum_{j=1}^n \exp(\text{sim}(q, p_j^-))}$$



Training with “in-batch” negatives

$$L(q, p^+, p_1^-, p_2^-, \dots, p_n^-) = -\log \frac{\exp(\text{sim}(q, p^+))}{\exp(\text{sim}(q, p^+)) + \sum_{j=1}^n \exp(\text{sim}(q, p_j^-))}$$

Back-propagation to all in-batch negatives!




Contriever (Izacard et al. 2022)

Unsupervised dense retrieval model!

- How to get **positive** samples in unlabeled text?
- How to use more **negative** samples during training?

Contriever (Izacard et al. 2022)


Independent Cropping



Positive pairs from unlabeled text!


Contriever (Izacard et al. 2022)

MoCo (Momentum Contrast)




Contriever (Izacard et al. 2022)

MoCo (Momentum Contrast)



Contriever (Izacard et al. 2022)

MoCo (Momentum Contrast)



Get more negatives from previous batches!

Retrieval-in-context in LM (Ram et al. 2023)

x = World Cup 2022 was the last with 32 teams, before the increase to

World Cup 2022 was the last with 32 teams, before the increase to



Retrieval Model



FIFA World Cup 2026 will expand to 48 teams. World Cup 2022 was the last with 32 teams, before the increase to



LM



48 in the 2026 tournament.

Retrieval-in-context in LM (Ram et al. 2023)

x = World Cup 2022 was the last with 32 teams, before the increase to

World Cup 2022 was the last with 32 teams, before the increase to



Retrieval Model

BM25, DPR, Contriever, ...



FIFA World Cup 2026 will expand to 48 teams. World Cup 2022 was the last with 32 teams, before the increase to




LM

GPT, OPT, LLaMA, ...



48 in the 2026 tournament.

Retrieval-in-context in LM




Better retrieval model



Better **retrieval-based LMs**

Better base LMs

Retrieval-in-context in LM



Better retrieval model



Better **retrieval-based LMs**

Better base LMs

Each component can be improved separately

kNN-LM (Khandelwal et al. 2020)

Inference

$$P_{k\text{NN}}(y \mid x) \propto \sum_{(k,v) \in \mathcal{D}} \mathbb{I}[v = y] \exp(-d(\text{Enc}(k), \text{Enc}(x)))$$

$$P_{k\text{NN-LM}}(y \mid x) = \lambda P_{\text{LM}}(y \mid x) + (1 - \lambda) P_{k\text{NN}}(y \mid x)$$

kNN-LM (Khandelwal et al. 2020)

Inference

Use the same encoder as the base LMs!

$$P_{k\text{NN}}(y \mid x) \propto \sum_{(k,v) \in \mathcal{D}} \mathbb{I}[v = y] \exp(-d(\text{Enc}(k), \text{Enc}(x)))$$

$$P_{k\text{NN-LM}}(y \mid x) = \lambda P_{\text{LM}}(y \mid x) + (1 - \lambda) P_{k\text{NN}}(y \mid x)$$

kNN-LM (Khandelwal et al. 2020)

Inference

Use the same encoder as the base LMs!


$$P_{k\text{NN}}(y|x) \propto \sum_{(k,v) \in \mathcal{D}} \mathbb{I}[v = y] \exp(-d(\text{Enc}(k), \text{Enc}(x)))$$

$$P_{k\text{NN-LM}}(y|x) = \lambda P_{\text{LM}}(y|x) + (1 - \lambda) P_{k\text{NN}}(y|x)$$

Training

Minimize $-\log P_{\text{LM}}(y|x)$

Independent training





Work with off-the-shelf models (no extra training required)



Each part can be improved independently

Independent training

-  Work with off-the-shelf models (no extra training required)
-  Each part can be improved independently
-  LMs are not trained to leverage retrieval
-  Retrieval models are not optimized for LM tasks/domains

Q & A



coffee break

We'll be back soon!


Section 4:

Retrieval-based LMs: Training (cont'd)

Why is training challenging?



Datastore



Training methods for retrieval-based LMs

- Independent training
- Sequential training
- Joint training w/ asynchronous index update
- Joint training w/ In-batch approximation

Training methods for retrieval-based LMs


- Independent training
- **Sequential training**
- Joint training w/ asynchronous index update
- Joint training w/ in-batch approximation

Sequential training

- One component is first trained independently and then fixed
- The other component is trained with an objective that depends on the first one


Sequential training

- One component is first trained independently and then fixed
- The other component is trained with an objective that depends on the first one




Sequential training

- One component is first trained independently and then fixed
- The other component is trained with an objective that depends on the first one




Sequential training

- One component is first trained independently and then fixed
- The other component is trained with an objective that depends on the first one



Sequential training

- Retrieval models are first trained independently and then fixed
- Language models are trained with an objective that depends on the retrieval



RETRO (Borgeaud et al. 2021)

x = World Cup 2022 was ~~the last with 32 teams,~~ before the increase to

x₁

x₂

x₃


RETRO (Borgeaud et al. 2021)

\mathbf{x} = World Cup 2022 was ~~the last with 32 teams~~, before the increase to

\mathbf{x}_1

\mathbf{x}_2

\mathbf{x}_3



RETRO (Borgeaud et al. 2021)


\mathbf{x} = World Cup 2022 was~~the~~ last with 32 teams,~~before~~ before the increase to

\mathbf{x}_1

\mathbf{x}_2

\mathbf{x}_3

(k chunks of text per split)



RETRO (Borgeaud et al. 2021)



\mathbf{x} = World Cup 2022 was~~the~~ last with 32 teams,~~before~~ the increase to

\mathbf{x}_1

\mathbf{x}_2



\mathbf{x}_3

(k chunks of text per split)





RETRO (Borgeaud et al. 2021)


\mathbf{x} = World Cup 2022 was~~the~~ the last with 32 teams,~~before~~ before the increase to




RETRO:Training




RETRO:Training



RETRO: Training




RETRO: Training




Sequential training

- One component is first trained independently and then fixed
- The other component is trained with an objective that depends on the first one




Sequential training


- Language models are first trained independently and then fixed
- Retrieval models are trained/fine-tuned with supervisions from LMs




REPLUG (Shi et al. 2023)




REPLUG (Shi et al. 2023)



REPLUG (Shi et al. 2023)



REPLUG LSR (LM-Supervised Retrieval)



Updating retrieval encoder → Retrieval Index becomes “stale”

How to deal with this issue? We will talk about it soon!

REPLUG results

Perplexity			
Model		# Parameters	Original
GPT-2	Small	117M	1.33
	Medium	345M	1.20
	Large	774M	1.19
	XL	1.5B	1.16
GPT-3 (black-box)	Ada	350M	1.05
	Babbage	1.3B	0.95
	Curie	6.7B	0.88
	Davinci	175B	0.80

REPLUG results

With Contriever, “**independent training**”

Model		# Parameters	Original	+ REPLUG	Gain %
GPT-2	Small	117M	1.33	1.26	5.3
	Medium	345M	1.20	1.14	5.0
	Large	774M	1.19	1.15	3.4
	XL	1.5B	1.16	1.09	6.0
GPT-3 (black-box)	Ada	350M	1.05	0.98	6.7
	Babbage	1.3B	0.95	0.90	5.3
	Curie	6.7B	0.88	0.85	3.4
	Davinci	175B	0.80	0.77	3.8

REPLUG results

Fine-tuning Contriever with
LM-supervised training
“Sequential training”

Model		# Parameters	Original	+ REPLUG	Gain %	+ REPLUG LSR	Gain %
GPT-2	Small	117M	1.33	1.26	5.3	1.21	9.0
	Medium	345M	1.20	1.14	5.0	1.11	7.5
	Large	774M	1.19	1.15	3.4	1.09	8.4
	XL	1.5B	1.16	1.09	6.0	1.07	7.8
GPT-3 (black-box)	Ada	350M	1.05	0.98	6.7	0.96	8.6
	Babbage	1.3B	0.95	0.90	5.3	0.88	7.4
	Curie	6.7B	0.88	0.85	3.4	0.82	6.8
	Davinci	175B	0.80	0.77	3.8	0.75	6.3

Sequential training

-  Work with off-the-shelf components (either a large index or a powerful LM)
-  LMs are trained to effectively leverage retrieval results
-  Retrievers are trained to provide text that helps LMs the most
-  One component is still fixed and not trained

Let's jointly train retrieval models and LMs!


Training methods for retrieval-based LMs

- Independent training
- Sequential training
- **Joint training w/ asynchronous index update**
- **Joint training w/ in-batch approximation**

Why is training challenging?




Datastore



Too large! Expensive to update index during training!

Challenges of updating retrieval models




Training methods for retrieval-based LMs


- Independent training
- Sequential training
- **Joint training w/ asynchronous index update**
- Joint training w/ in-batch approximation

Asynchronous index update

- Retrieval models and language models are trained jointly
- Allow the index to be “**stale**”; rebuild the retrieval index every T steps




Asynchronous index update



Asynchronous index update




Datastore



Asynchronous index update




Datastore




REALM (Guu et al. 2020)

x = The [MASK] at the top of the pyramid.




REALM: Training

Objective: maximize $\sum_{z \in \mathcal{Z}_\theta} P_\theta(z | q) P_\theta(y | q, z)$




REALM: Training

Objective: maximize $\sum_{z \in \mathcal{Z}_\theta} P_\theta(z | q)P_\theta(y | q, z)$



REALM: Training

Objective: maximize $\sum_{z \in \mathcal{Z}_\theta} P_\theta(z | q)P_\theta(y | q, z)$



REALM: Index update rate

How often should we update the retrieval index?

- Too often: expensive
- Too slowly: out-dated

REALM: Index update rate

How often should we update the retrieval index?

- Too often: expensive
- Too slowly: out-dated


REALM: updating the index every 500 training steps

REALM: Index update rate


How often should we update the retrieval index?

- Too often: expensive
- Too slowly: out-dated

REALM: updating the index every 500 training steps




Atlas (Izacard et al. 2022)




Atlas (Izacard et al. 2022)

Retrieval-based generative seq2seq model




Atlas (Izacard et al. 2022)




Process each doc independently using “Fusion-in-Decoder”

Atlas (Izacard et al. 2022)

Jobs is CEO of _



Atlas (Izacard et al. 2022)




Adapted to a lot of downstream tasks! (Section 5)

Atlas: Retriever training

Perplexity Distillation

Retrieve the text that can help LM encoders improve perplexity

$$P_{\text{retr}}(z | q) = \frac{\exp(s(z, q))}{\sum_{k=1}^K \exp(s(z_k, q))}$$



$$P_{\text{ppl}}(z | q, y) = \frac{\exp(\log P_{\text{LM}}(y | q, z))}{\sum_{k=1}^K \exp(\log P_{\text{LM}}(y | q, z_k))}$$

How likely each document is retrieved

How much each document improves the ppl

Atlas: Retriever training

Similarity based on retrieval encoder

$$P_{\text{retr}}(z | q) = \frac{\exp(s(z, q))}{\sum_{k=1}^K \exp(s(z_k, q))}$$

How likely each document is retrieved

KL Divergence


Prob of the gold labels if augmenting this text chunk

$$P_{\text{ppl}}(z | q, y) = \frac{\exp(\log P_{\text{LM}}(y | q, z))}{\sum_{k=1}^K \exp(\log P_{\text{LM}}(y | q, z_k))}$$

How much each document improves the ppl


Perplexity Distillation

Atlas:Asynchronous index update



Update the index every T steps
30% overhead for asynchronous updating

Atlas: Asynchronous index update



Update the index every T steps


30% overhead for asynchronous updating

How can we get rid of this?

Training methods for retrieval-based LMs

- Independent training
- Sequential training
- Joint training w/ asynchronous index update
- **Joint training w/ in-batch approximation**

In-batch approximation



TRIME:Training with in-batch memory (Zhong et al. 2022)

Similar to kNN-LM

Datastore $\mathcal{D} = \{(k, v)\}$

context next token

Keys	Values
<i>10/10, would buy this</i>	<i>cheap</i>
<i>Item delivered broken. Very</i>	<i>cheap</i>
<i>To check the version of PyTorch, you can use</i>	<i>torch</i>
<i>You are permitted to bring a</i>	<i>torch</i>
<i>A group of infections ... one of the</i>	<i>torch</i>

TRIME: Training with in-batch memory (Zhong et al. 2022)

Similar to kNN-LM

Datastore $\mathcal{D} = \{(k, v)\}$

context next token

Keys	Values
<i>10/10, would buy this</i>	<i>cheap</i>
<i>Item delivered broken. Very</i>	<i>cheap</i>
<i>To check the version of PyTorch, you can use</i>	<i>torch</i>
<i>You are permitted to bring a</i>	<i>torch</i>
<i>A group of infections ... one of the</i>	<i>torch</i>

Inference

$$P(y | x) \propto \exp(E^\top f(x)) + \sum_{(k,v) \in \mathcal{D}} \mathbb{I}[v = y] \exp(-d(\text{Enc}(k), \text{Enc}(x)))$$

output embedding
(same as standard LMs)

datastore
(very large!)

TRIME: Training with in-batch memory (Zhong et al. 2022)

Similar to kNN-LM

Datastore $\mathcal{D} = \{(k, v)\}$

context next token

Keys	Values
10/10, would buy this	cheap
Item delivered broken. Very	cheap
To check the version of PyTorch, you can use	torch
You are permitted to bring a	torch
A group of infections ... one of the	torch

Inference

$$P(y | x) \propto \exp(E^\top f(x)) + \sum_{(k,v) \in \mathcal{D}} \mathbb{I}[v = y] \exp(-d(\text{Enc}(k), \text{Enc}(x)))$$

output embedding
(same as standard LMs)

1. Aligning the output representations with static embeddings


datastore
(very large!)

2. Aligning input context with all the contexts in datastore that share the same next token

TRIME: Training

Key idea: build a temporary index from same training batch on the fly

$$P(y \mid x) \propto \exp(E^\top f(x)) + \sum_{(k,v) \in \mathcal{D}_{train}} \mathbb{I}[v = y] (-d(\text{Enc}(k), \text{Enc}(x)))$$



In-batch approximation

(built from in-batch examples on the fly)

TRIME: Training

Key idea: build a temporary index from same training batch on the fly

$$P(y | x) \propto \exp(E^\top f(x)) + \sum_{(k,v) \in \mathcal{D}_{train}} \mathbb{I}[v = y] (-d(\text{Enc}(k), \text{Enc}(x)))$$

In-batch approximation

(built from in-batch examples on the fly)

We can back-propagate to all the representations in datastore \mathcal{D}_{train} !

TRIME: full index vs. in-batch index



Full corpus



Keys	Values
<i>To check the version of PyTorch, you can use</i>	<i>torch</i>
<i>Item delivered broken. Very</i>	<i>cheap</i>
<i>He moves to</i>	<i>Apple</i>
<i>Apple merged with NeXT, and Jobs became</i>	<i>CEO</i>
<i>...</i>	<i>...</i>

Full index (used during inference)

TRIME: full index vs. in-batch index



Full corpus



Keys	Values
To check the version of PyTorch, you can use	torch
Item delivered broken. Very	cheap
He moves to	Apple
Apple merged with NeXT, and Jobs became	CEO
...	...

Full index (used during inference)

Compute on the fly!

Apple merged with NeXT, and ...
VS Code was developed by Microsoft for Windows in 2015 ...
He moves to Apple ...
...




Keys	Values
Jobs	became
Apple merged with NeXT, and Jobs became	CEO
...	...
He moves to	Apple
VS Code was developed	By

Training batch

In-batch index (used during training)


TRIME: full index vs. in-batch index

How to batch training data —
so we can have good in-batch examples?



TRIME: Data batching strategy


■ Current token ■ In memory ■ Not in memory



Key idea: **similar contexts** — more training signals from in-batch examples!


Use **BM25** scores to find similar contexts

TRIME: Results




Model size: 247M
Sequence len: 3072

NPM: Nonparametric masked LMs (Min et al. 2023)




NPM: Nonparametric masked LMs (Min et al. 2023)




1. **masked** language model pertained on >1B tokens

NPM: Nonparametric masked LMs (Min et al. 2023)



NPM: Nonparametric masked LMs (Min et al. 2023)



1. masked language model pertained on >1B tokens
2. Each mask corresponds to a phrase (instead of a token)
3. During inference, predictions are made **purely** according to retrieval results

NPM: Nonparametric masked LMs (Min et al. 2023)

Key challenge: how to train the retrieval-based model on >1B tokens?

1. How to approximate full corpus retrieval
2. How to choose positives and negatives

In-batch approximation with same-doc batching

NPM:Training

1. Sample sequences from the same document

The 2010 Seattle Seahawks Were a Playoff Team With a Losing Record

by Julie Rhoads in More Articles: NFL
Published on January 9, 2021 | View Comments

SHARE:

Making the NFL playoffs is tough. Only the best of the best make it to the postseason, or do they? Throughout the years, [plenty of subpar teams](#) with losing records have made it through. Whether it's because of a fluke play, bad calls, or just plain luck, there's always one team that seems a bit out of its league once the playoffs begin. The [Seattle Seahawks](#) were that team in 2010.

Heading to the playoffs with a losing record

The 2010 Seattle Seahawks were part of a losing division that year. It was [Pete Carroll's first season](#) as head coach and quarterback Matt Hasselbeck's last with the team. The season started off strong with Seattle having a 4-2 record, but things only went downhill from there. Injuries and poor play caused the Seahawks to lose seven of their last 10 games. They ended the season 7-9 but still won the division and a trip to the playoffs since the other divisional teams had worse records.

In the 2010 NFL season, the Seattle Seahawks made history by making it into the playoffs despite having a 7-9 record.

... against the Seattle Seahawks as a member of (...) In the 2010 season, the Seahawks became the first team in NFL history to ...

For simplicity, we assume 2 sequences in a batch

NPM:Training

2. Identify co-occurring spans

The 2010 Seattle Seahawks Were a Playoff Team With a Losing Record

by Julie Rhoads [in](#) | More Articles: NFL
Published on January 9, 2021 | [View Comments](#)

SHARE:

Making the NFL playoffs is tough. Only the best of the best make it to the postseason, or do they? Throughout the years, [plenty of subpar teams](#) with losing records have made it through. Whether it's because of a fluke play, bad calls, or just plain luck, there's always one team that seems a bit out of its league once the playoffs begin. The [Seattle Seahawks](#) were that team in 2010.

Heading to the playoffs with a losing record

The 2010 Seattle Seahawks were part of a losing division that year. It was [Pete Carroll's first season](#) as head coach and quarterback Matt Hasselbeck's last with the team. The season started off strong with Seattle having a 4-2 record, but things only went downhill from there. Injuries and poor play caused the Seahawks to lose seven of their last 10 games. They ended the season 7-9 but still won the division and a trip to the playoffs since the other divisional teams had worse records.

In the 2010 NFL season, **the Seattle Seahawks** made history by making it into the playoffs despite having a 7-9 record.

... against **the Seattle Seahawks** as a member of (...) In the 2010 season, the Seahawks became the first team in NFL history to ...

For simplicity, we assume 2 sequences in a batch

NPM:Training

3. One is “positive” for the other

The 2010 Seattle Seahawks Were a Playoff Team With a Losing Record

by Julie Rhoads [in](#) | More Articles: NFL
Published on January 9, 2021 | [View Comments](#)

SHARE:

Making the NFL playoffs is tough. Only the best of the best make it to the postseason, or do they? Throughout the years, [plenty of subpar teams](#) with losing records have made it through. Whether it's because of a fluke play, bad calls, or just plain luck, there's always one team that seems a bit out of its league once the playoffs begin. The [Seattle Seahawks](#) were that team in 2010.

Heading to the playoffs with a losing record

The 2010 Seattle Seahawks were part of a losing division that year. It was [Pete Carroll's first season](#) as head coach and quarterback Matt Hasselbeck's last with the team. The season started off strong with Seattle having a 4-2 record, but things only went downhill from there. Injuries and poor play caused the Seahawks to lose seven of their last 10 games. They ended the season 7-9 but still won the division and a trip to the playoffs since the other divisional teams had worse records.

In the 2010 NFL season, _____ made history by making it into the playoffs despite having a 7-9 record.

... against **the Seattle Seahawks** as a member of (...) In the 2010 season, the Seahawks became the first team in NFL history to ...

For simplicity, we assume 2 sequences in a batch

NPM:Training

4. The others are “negatives”

The 2010 Seattle Seahawks Were a Playoff Team With a Losing Record

by Julie Rhoads in More Articles: NFL
Published on January 9, 2021 | View Comments

SHARE:   

Making the NFL playoffs is tough. Only the best of the best make it to the postseason, or do they? Throughout the years, [plenty of subpar teams](#) with losing records have made it through. Whether it's because of a fluke play, bad calls, or just plain luck, there's always one team that seems a bit out of its league once the playoffs begin. The [Seattle Seahawks](#) were that team in 2010.

Heading to the playoffs with a losing record

The 2010 Seattle Seahawks were part of a losing division that year. It was [Pete Carroll's first season](#) as head coach and quarterback Matt Hasselbeck's last with the team. The season started off strong with Seattle having a 4-2 record, but things only went downhill from there. Injuries and poor play caused the Seahawks to lose seven of their last 10 games. They ended the season 7-9 but still won the division and a trip to the playoffs since the other divisional teams had worse records.

In the 2010 NFL season, _____ made history by making it into the playoffs despite having a 7-9 record.

positive

... against **the Seattle Seahawks** as a member of (...) In the 2010 season, the Seahawks became the first team in NFL history to ...

For simplicity, we assume 2 sequences in a batch

Beyond lexical clues?

In TRIME and NPM, retrieval models are trained to use *lexical* information

Positives: **co-occurring** tokens/spans

Can we do more than that?

RPT: Retrieval-pretrained transformer

(Rubin and Berant 2023)

Reference score $P(\text{"Apple"} \mid \text{"Jobs become CEO of"}, \underline{\text{"NeXT merged with ..."}})$

Reference chunk

RPT: Retrieval-pretrained transformer

(Rubin and Berant 2023)

Reference score $P(\text{"Apple"} \mid \text{"Jobs become CEO of"}, \text{"NeXT merged with ..."})$

Reference chunk

$P(\text{"Apple"} \mid \text{"Jobs become CEO of"}, \text{"He joined his former ..."}) > \text{Reference score}$

Positive chunks

RPT: Retrieval-pretrained transformer

(Rubin and Berant 2023)

Reference score $P(\text{"Apple"} \mid \text{"Jobs become CEO of"}, \text{"NeXT merged with ..."})$

Reference chunk


$P(\text{"Apple"} \mid \text{"Jobs become CEO of"}, \text{"He joined his former ..."}) > \text{Reference score}$

Positive chunks

$P(\text{"Apple"} \mid \text{"Jobs become CEO of"}, \text{"Jobs was raised ..."}) < \text{Reference score}$

Negative chunks

Joint training

-  End-to-end trained – each component is optimized
-  Good performance
-  Training is more complicated
-  Train-test discrepancy still remains

Summary

Training method	+	-
Independent training (Ram et al 2023; Khandelwal et al 2020)		
Sequential training (Borgeaud et al 2021; Shi et al 2023)	<ul style="list-style-type: none">* Easy to implement: off-the-shelf models* Easy to improve: sub-module can be separately improved	<ul style="list-style-type: none">* Models are not end-to-end trained – suboptimal performance
Joint training: async update (Guu et al 2020; Izacard et al 2022)	<ul style="list-style-type: none">* End-to-end trained – very good performance!	<ul style="list-style-type: none">* Training may be complicated (overhead, batching methods, etc)* Train-test discrepancy still remains
Joint training: in-batch approx (Zhong et al 2022; Min et al 2023; Rubin and Berant 2023)		

How do retrieval-based language models perform on downstream tasks? → **Section 5!**